

## Java Concurrency Utilities

Based on JavaOne talk given by  
**David Holmes & Brian Goetz**

### Overview

- Rationale and goals for JSR 166
  - Java community process – concurrency utilities
- Executors – thread pools and scheduling
- Futures
- Concurrent Collections
- Locks, conditions and synchronizers
- Atomic variables

## Why Concurrency Utilities

- Java's built-in concurrency primitives – wait(), notify(), and synchronized:
  - Hard to use correctly
  - Easy to use incorrectly
  - Too low level for many applications
  - Can lead to poor performance if used incorrectly
  - Leave out lots of useful concurrency constructs

## Goals

- Provide efficient, correct & reusable concurrency building blocks
- Enhance scalability, performance, readability, maintainability, and thread-safety of concurrent Java applications

## Background

- Found in `java.util.concurrent`
  - based on Doug Lea's `EDU.oswego.cs.dl.util.concurrent` package
- APIs take advantage of native JVM constructs & Java Memory Model guarantees specified in JSR 133

## Building Blocks

- Executor, ThreadPool, and Future
- Concurrent collections:
  - `BlockingQueue`, `ConcurrentHashMap`, `CopyOnWriteArray`
- Locks and Conditions
- Synchronizers: Semaphores, Barriers, etc.
- Atomic Variables
  - Low-level compare-and-set operation

## Executor

- Standardizes asynchronous invocation
- Separates job submission from execution policy
  - `anExecutor.execute(aRunnable)`
  - not `new Thread(aRunnable).start()`
- Two code styles supported:
  - Actions: **Runnable**s
  - Functions: **Callable**s
  - Also has lifecycle mgmt: e.g., cancellation, shutdown
- Executor usually created via **Executors** factory class
  - Configures **ThreadPoolExecutor**
  - Customizes shutdown methods, before/after hooks, saturation policies, queuing

## Executor & ExecutorService

- ExecutorService adds lifecycle management to Executor

```
public interface Executor {
    void execute(Runnable command);
}
public interface ExecutorService extends Executor {
    void shutdown();
    List<Runnable> shutdownNow();
    boolean isShutdown();
    boolean isTerminated();
    boolean awaitTermination( long timeout, TimeUnit unit);
    // other convenience methods for submitting tasks
}
```

## Creating Executors

- Executors factory methods

```
public class Executors {
    static ExecutorService newSingleThreadedExecutor();
    static ExecutorService newFixedThreadPool(int n);
    static ExecutorService newCachedThreadPool(int n);
    static ScheduledExecutorService newScheduledThreadPool(int n);
    // additional versions & utility methods
}
```

## (Not) Executor Example

- Thread per message Web Server (no limit on thread creation)

```
class WebServer {
    public static void main( String [] args) {
        ServerSocket socket = new ServerSocket ( 80 );
        while ( true ) {
            final Socket connection = socket.accept();
            Runnable r = new Runnable () {
                public void run () {handleRequest(connection);}
            };
            new Thread (r).start();
        }
    }
}
```

## Executor Example

- Thread pool web server - better resource management
- **class** WebServer {  
**Executor** pool = Executors.newFixedThreadPool(7);  
**public static void** main(**String**[] args) {  
**ServerSocket** socket = new **ServerSocket**(80);  
**while** (**true**) {  
final **Socket** connection = socket.accept();  
**Runnable** r = new **Runnable**() {  
**public void** run() {handleRequest(connection);}  
};  
pool.execute(r);  
}  
}  
}

## Future and Callable

- Callable is functional analog of Runnable  
**interface** Callable<**V**> {  
**V** call() **throws** **Exception**;  
}
- Future represents result of asynchronous computation  
**interface** Future<**V**> {  
**V** get() **throws** **InterruptedException**, **ExecutionException**;  
**V** get(**long** timeout, **TimeUnit** unit);  
**boolean** cancel(**boolean** mayInterrupt);  
**boolean** isCancelled();  
**boolean** isDone();  
}

## Using Futures

- Client initiates asynchronous computation via oneway message
- Client receives a “handle” to the result: a Future
- Client does other work while waiting for result
- When ready Client requests result from Future, blocking if necessary until result is available
- Client uses result

## FutureTask

- A cancellable asynchronous computation
- A base implementation of Future
- Can wrap a Callable or Runnable
  - Allows FutureTask to be submitted to an Executor

## Future Example

- See: FutureStringReverser.java
- See: FutureTaskStringReverser.java

## Another Future Example

- Implementing a cache with Future

```
public class Cache<K, V>
{
    Map<K, Future<V>> map = new ConcurrentHashMap();
    Executor executor = Executors.newFixedThreadPool(8);
    public V get (final K key) {
        Future<V> f = map.get(key); // null if key not found
        if (f == null) { Callable<V> c = new Callable<V>() {
            public V call() { // compute value associated with key
            };
            f = new FutureTask<V>(c);
            Future old = map.putIfAbsent(key, f); // if key not found put(key,f) & return null
            if (old == null) { // otherwise return get(key)
                executor.execute(f);
            } else { f = old; }
            return f.get();
        }
    }
}
```



## ScheduledExecutorService

- For deferred and recurring tasks, can schedule
  - **Callable or Runnable to run once** with a fixed delay after submission
  - Schedule a **Runnable to run periodically at a fixed rate**
  - Schedule a **Runnable to run periodically with a fixed delay** between executions
- Submission returns a **ScheduledFutureTask** handle which can be used to cancel the task
- Like **Timer**, but supports pooling and is more robust

## Concurrent Collections

- Pre-1.5 Java class libraries had few concurrent (vs. synchronized) classes
  - Synchronized collections:
    - Hashtable, Vector, and Collections.synchronized\*
    - Often required locking during iteration
    - Locking becomes is a source of contention
- Java 1.5 concurrent collections:
  - Allow multiple operations to overlap
    - Some differences in semantics

## Queues

- Queue interface added to java.util

```
interface Queue<E> extends Collection<E> {
    boolean offer(E x); // try to insert
    E poll(); //retrieve and remove. Return null if empty
    E remove() throws NoSuchElementException; //retrieve and remove
    E peek(); // retrieve, don't remove. Return null if empty
    E element() throws NoSuchElementException; //retrieve, don't remove
}
```

- Thread-safe and non-thread safe implementations
  - Non-thread-safe - **LinkedList**
  - Non-thread-safe - **PriorityQueue**
  - Thread-safe non-blocking - **ConcurrentLinkedQueue**

## Blocking Queues

- Extends **Queue** to provide blocking operations
  - Retrieval: wait for queue to become nonempty
  - Insertion: wait for capacity to be available
- Common in producer-consumer designs
- Can be bounded or unbounded
- Implementations provided:
  - **LinkedBlockingQueue** (FIFO, may be bounded)
  - **PriorityBlockingQueue** (priority, unbounded)
  - **ArrayBlockingQueue** (FIFO, bounded)
  - **SynchronousQueue** (rendezvous channel)
- See API for details

## Producer-Consumer Examples

- See:
  - `ProducerConsumerPrimitive.java` (wait/notify)
  - `ProducerConsumerConcUtil.java` (BlockingQueue)

## Concurrent Collections

- **`ConcurrentHashMap`** - Concurrent (scalable) alternative to **`Hashtable`** or **`Collections.synchronizedMap`**
  - Multiple reads can overlap each other
  - Reads can overlap writes
  - Retrieval operations reflect the results of the most recently completed update operations holding at onset of operation
  - Up to 16 writes can overlap
  - Iterators do not throw **`ConcurrentModificationException`**
- **`CopyOnWriteArrayList`**
  - Optimized for case where iteration is much more frequent than insertion or removal. E.g., event listeners

## Performance Comparison

- **ConcurrentHashMap vs. Collections.synchronizedMap**
- See HashMapPerfTest.java
- **Note:** incrementCount() is not safe

## Locks and Lock Support

- High-level locking interface
- Adds non-blocking lock acquisition

```
interface Lock {  
    void lock();  
    void lockInterruptibly() throws IE;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit) throws IE;  
    void unlock();  
    Condition newCondition() throws UnsupportedOperationException;  
}
```

## ReentrantLock

- Flexible, high-performance lock implementation
- Implements a reentrant mutual exclusion lock (like Java intrinsic locks) but with extra features
  - Can interrupt a thread waiting to acquire a lock
  - Can specify a timeout while waiting for a lock
  - Can poll for lock availability
  - Can have multiple wait-sets per lock via the **Condition** interface
- Outperforms built-in monitor locks in most cases, but slightly less convenient to use (requires finally block to release lock)

## Lock Example

- Locks not automatically released
  - Must release lock in **finally** block

```
Lock lock = new ReentrantLock();
...
lock.lock();
try {
    // perform operations protected by lock
} catch (Exception ex) {
    // restore invariants
} finally {
    lock.unlock();
}
```

## ReadWrite Locks

- **ReadWriteLock** interface defines a pair of locks;
  - one for readers; one for writers

```
interface ReadWriteLock {
    Lock readLock();
    Lock writeLock();
}
```

- **ReentrantReadWriteLock** class
  - Multiple readers, single writer
  - Allows writer to acquire read lock
  - Allows writer to downgrade to read lock
  - Supports “fair” and “non-fair” (default) acquisition

## Read/Write Lock Example

```
class RWDictionaryRWL {
    private final Map<String, Data> m = new TreeMap<String, Data>();
    private final ReentrantReadWriteLock rwl = new ReentrantReadWriteLock();
    private final Lock r = rwl.readLock();
    private final Lock w = rwl.writeLock();
    public Data get(String key) {
        r.lock();
        try { return m.get(key); }
        finally { r.unlock(); }
    }
    public Data put(String key, Data value)
    {
        w.lock();
        try { return m.put(key, value); }
        finally { w.unlock(); }
    }
}
```

## Read/Write Lock Example

- See
  - [RWDictionary.java](#) & [RWDictionaryRWL.java](#)

## Condition

- Condition lets you wait for a condition to hold (like wait), but adds several features

```
interface Condition {  
    void await() throws IE;  
    boolean await( long time, TimeUnit unit ) throws IE;  
    long awaitNanos( long nanosTimeout ) throws IE;  
    void awaitUninterruptibly()  
    boolean awaitUntil( Date deadline ) throws IE;  
    void signal();  
    void signalAll();  
}
```

## Condition (cont.)

- Many improvements over wait()/notify()
  - Multiple conditions per lock
  - Absolute and relative time-outs
  - Timed waits tell you why you returned
  - Convenient uninterruptible wait

## Condition Example

```

class BoundedBufferCond {
    Lock lock = new ReentrantLock ();
    Condition notFull = lock.newCondition();
    Condition notEmpty = lock.newCondition();
    Object[] items = new Object[100];
    int putptr, takeptr, count;

    public void put( Object x) throws IE {
        lock.lock();
        try {
            while (count == items.length) notFull.await();
            items[putptr] = x;
            if (++putptr == items.length) putptr = 0;
            ++count;
            notEmpty.signal();
        } finally { lock.unlock(); }
    }
}

```



### Condition Example (cont.)

```
public Object take() throws IE {  
    lock.lock();  
    try {  
        while (count == 0) notEmpty.await();  
        Object x = items[takeptr];  
        if (++takeptr == items.length) takeptr = 0;  
        --count;  
        notFull.signal();  
        return x;  
    } finally { lock.unlock(); }  
}
```

### Condition Example (cont.)

- Previous example in BoundedBufferCond.java
- See also: BoundedBufferPrim.java

## Synchronizers

- Utilities for coordinating access and control
- **CountDownLatch** – allows one or more threads to wait for a set of threads to complete an action
- **CyclicBarrier** – allows a set of threads to wait until they all reach a specified barrier point
- **Semaphore** – Dijkstra counting semaphore, managing some number of permits
- **Exchanger** – allows two threads to rendezvous and exchange data, such as exchanging an empty buffer for a full one

## CountDownLatch

- Latching variables are conditions that once set never change
- Often used to start several threads, but have them wait for a signal before continuing
- See: `CountDownLatchTest.java`

## CyclicBarrier

- Allows threads to wait at a common barrier point
- Useful when a fixed-sized party of threads must occasionally wait for each other
- Cyclic Barriers can be re-used after threads released
- Can execute a Runnable once per barrier point
  - After the last thread arrives, but before any are released
  - Useful for updating shared-state before threads continue
- See: CyclicBarrierEx1.java & CyclicBarrierEx2.java

## Semaphore

- Semaphore maintain a logical set of permits
- acquire() blocks until a permit is free, then takes it
- release() adds a permit, releasing a blocking acquirer
- Often used to restrict the number of threads that can access some resource
  - But can be used to implement many sync disciplines
- See: SemaphoreTunnel.java & SemaphoreBuffer.java

## Exchanger

- Synch. point where two threads exchange objects
- A bidirectional SynchronizedQueue
- Each thread presents some object on entry to the exchange() method, and receives the object presented by the other thread on return
- See ExchangerTest.java

## Atomic Variables

- Holder classes for scalars, references and fields
- Supports atomic operations
  - Compare-and-set (CAS)
  - Get and set and arithmetic (where applicable)
- Ten main classes: { int, long, ref } X { value, field, array }
  - E.g. **AtomicInteger** useful for counters, sequences, statistics
- Essential for writing efficient code on MPs
  - Nonblocking data structures & optimistic algorithms
  - Reduce overhead/contention updating “hot” fields
- JVM uses best construct available on platform
  - CAS, load-linked/store-conditional, locks

## Atomic Variables

- See: CounterTest.java